

Multiple Linear Regression with Dummies - Exercise

You are given a real estate dataset.

Real estate is one of those examples that every regression course goes through as it is extremely easy to understand and there is a (almost always) certain causal relationship to be found.

The data is located in the file: 'real_estate_price_size_year_view.csv'.

You are expected to create a multiple linear regression (similar to the one in the lecture), using the new data.

In this exercise, the dependent variable is 'price', while the independent variables are 'size', 'year', and 'view'.

Regarding the 'view' variable:

There are two options: 'Sea view' and 'No sea view'. You are expected to create a dummy variable for view and include it in the regression

Good luck!

Import the relevant libraries

```
In [4]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import statsmodels.api as sm
import seaborn as sn
sn.set()
```

Load the data

```
In [5]: data = pd.read_csv('real_estate_price_size_year_view.csv')
```

```
In [34]: data.head()
```

```
Out[34]:
```

	price	size	year	view
0	234314.144	643.09	2015	No sea view
1	228581.528	656.22	2009	No sea view
2	281626.336	487.29	2018	Sea view
3	401255.608	1504.75	2015	No sea view
4	458674.256	1275.46	2009	Sea view

```
In [38]: data.describe(include='all')
```

```
Out[38]:
```

	price	size	year	view
count	100.000000	100.000000	100.000000	100
unique	NaN	NaN	NaN	2
top	NaN	NaN	NaN	No sea view
freq	NaN	NaN	NaN	51
mean	292289.470160	853.024200	2012.600000	NaN
std	77051.727525	297.941951	4.729021	NaN
min	154282.128000	479.750000	2006.000000	NaN
25%	234280.148000	643.330000	2009.000000	NaN
50%	280590.716000	696.405000	2015.000000	NaN
75%	335723.696000	1029.322500	2018.000000	NaN
max	500681.128000	1842.510000	2018.000000	NaN

```
In [35]: data_copy = data.copy()
```

Create a dummy variable for 'view'

```
In [41]: data_copy['view'] = data_copy['view'].map({'Sea view':1, 'No sea view':0})
```

```
In [42]: data_copy
```

```
Out[42]:
```

	price	size	year	view
0	234314.144	643.09	2015	0
1	228581.528	656.22	2009	0
2	281626.336	487.29	2018	1
3	401255.608	1504.75	2015	0
4	458674.256	1275.46	2009	1
...
95	252460.400	549.80	2009	1
96	310522.592	1037.44	2009	0
97	383635.568	1504.75	2006	0
98	225145.248	648.29	2015	0
99	274922.856	705.29	2006	1

100 rows × 4 columns

Create the regression

Declare the dependent and the independent variables

```
In [43]: y = data_copy['price']  
x1 = data_copy[['size', 'year', 'view']]
```

Regression

```
In [44]: x = sm.add_constant(x1)
results = sm.OLS(y,x).fit()
results.summary()
```

Out[44]: OLS Regression Results

Dep. Variable:	price	R-squared:	0.913
Model:	OLS	Adj. R-squared:	0.910
Method:	Least Squares	F-statistic:	335.2
Date:	Wed, 25 Aug 2021	Prob (F-statistic):	1.02e-50
Time:	23:51:01	Log-Likelihood:	-1144.6
No. Observations:	100	AIC:	2297.
Df Residuals:	96	BIC:	2308.
Df Model:	3		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	-5.398e+06	9.94e+05	-5.431	0.000	-7.37e+06	-3.43e+06
size	223.0316	7.838	28.455	0.000	207.473	238.590
year	2718.9489	493.502	5.510	0.000	1739.356	3698.542
view	5.673e+04	4627.695	12.258	0.000	4.75e+04	6.59e+04

Omnibus:	29.224	Durbin-Watson:	1.965
Prob(Omnibus):	0.000	Jarque-Bera (JB):	64.957
Skew:	1.088	Prob(JB):	7.85e-15
Kurtosis:	6.295	Cond. No.	9.42e+05

Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 9.42e+05. This might indicate that there are strong multicollinearity or other numerical problems.

In []: